

Demo-Notes

Friday, September 17, 2021 7:35 PM

To learn some basic SQL, I decided to create a local test environment to demonstrate a Java application reading data from a Microsoft SQL server, transforming it, and then writing it to a Postgres Database.

This is a reasonable simulation, but I realize that it is missing many of the aspects needed for actual production use. Examples of a few areas where I simplified for the demo that would need to be addressed include:

- 1) The transformations are simple and would be much more complex in production.
- 2) Better security would be needed such as stronger passwords and keeping passwords in an external vault and not in SQL files and in Java code files
- 3) Use connection pooling to the databases

Technologies used to create the demo:

- VMware Virtual Machine
- Ubuntu 20.04 running in a VM
- Docker Engine for Linux
 - Postgres v12.8 image pulled from DockerHub and run in a Docker Container
 - Microsoft SQL 2019 image pulled from DockerHub and run in a Docker Container
- psql CLI
- sqlcmd CLI
- Dev stack:
 - Java - openJDK 11
 - java.sql - Java client library that supports ODBC connections to Postgres and to Microsoft SQL Server
 - IntelliJ IDE
- Version Control System: git
- Build System: Gradle
 - Maven Repository - For finding and pulling Java Library dependencies in Gradle - <https://mvnrepository.com/>
- Java source files
 - Main.java
 - Pg.java - contains methods for inserting and selecting from the Postgres database
 - MsSql.java - contains methods for inserting and selecting from the Microsoft SQL Server database
- T-SQL - Microsoft Transact SQL
- Postgres SQL

Detailed description of how I built the demo environment

- 1) Create a Virtual Machine on VMware
 - a. Install Linux (Ubuntu 20.04) Operating system in the VM
- 2) Setup the version control environment (git)
 - a. Add a repo on Github and git clone it to a local repo
 - b. Set up SSH key on Linux to allow git to authenticate with an SSH public key installed on Github
 - eval "\$(ssh-agent -s)"
 - ssh-add ~/.ssh/id_ed25519
 - c. Set files to not push to the remote git repo by editing .gitignore

```
# IntelliJ J IDE files
**/.idea/

# Common Ignore Folders
**/build/
**/.gradle/
```

- 3) Install Docker Engine on Linux - <https://docs.docker.com/engine/install/ubuntu/>

```

sudo apt-get install \apt-transport-https \ca-certificates \curl \gnupg \lsb-release

curl -fsSL https://download.docker.com/linux/ubuntu/gpg | sudo gpg --dearmor-
o/usr/share/keyrings/docker-archive-keyring.gpg

echo \"deb [arch=amd64 signed-by=/usr/share/keyrings/docker-archive-keyring.gpg]
https://download.docker.com/linux/ubuntu \
$(lsb_release -cs) stable\" | sudo tee /etc/apt/sources.list.d/docker.list >/dev/null

sudo apt-get install docker-ce docker-ce-cli containerd.io

```

- 4) Download the Postgres Container image from Docker Hub
 - a. Postgres Version 12.8
`sudo docker pull postgres:12.8`
 - b. To view local docker images: `docker image ls`
- 5) Run the Postgres Docker Container
 - a. To run a docker image to start the docker container (version 12.8 Postgres Docker image name is postgres:12.8):
`sudo docker run --name pg -p 5432:5432 -e POSTGRES_PASSWORD=pw -d postgres:12.8`
 - a. Docker command line: <https://docs.docker.com/engine/reference/commandline/cli/>
 - 1) Quick reference:
<https://www.docker.com/sites/default/files/d8/2019-09/docker-cheat-sheet.pdf>
 - 2) To stop a container:
`sudo docker container stop pg`
 - 3) To restart a stopped container:
`sudo docker container restart pg`
 - 4) To delete a container
`sudo docker container rm pg`
 - 5) To view the logs for a container
`sudo docker logs pg`
- 6) Install the Postgres command-line interface (CLI) - psql

<https://stackoverflow.com/questions/28290488/get-error-you-must-install-at-least-one-postgresql-client-version-package-whe>
- 7) Setup a database and a table in the Postgres server running in the Docker Container
 - 1) Connect to the Postgres SQL Server running in the docker container using an interactive psql session
`psql -h localhost -p 5432 -U postgres`
 - a. User is postgres
 - b. PASSWORD=pw
 - 2) Show the postgres version: `psql --version`
 - 3) Edit a pg_demo.sql file that will create a database and a table in that database in the Postgres server
 - a. Create a Postgres database - create database pg_testdb
https://www.tutorialspoint.com/postgresql/postgresql_create_database.htm
 - 1) Verify doing an interactive psql session: \l
 - b. To prevent an error, you can use the following to check for the database existence before creating it

`SELECT 'CREATE DATABASE <your db name>' WHERE NOT EXISTS (SELECT FROM pg_database WHERE datname = '<your db name>')\gexec`
 - c. Create the airplane table in Postgres -
https://www.tutorialspoint.com/postgresql/postgresql_create_table.htm
 - 1) To skip creating the table if it already exists, add "IF NOT EXISTS" as in the command below

`CREATE TABLE IF NOT EXISTS table_name (`
 - d. To show the tables in a database

`\c database_name`

```
\d
SELECT * FROM pg_testdb.*;
```

- e. To show the rows in a table

```
SELECT * FROM table_name;
```

- f. To insert a row into a table in Postgres

```
INSERT INTO table_name(column1, column2, ...) VALUES (value1, value2, ...);
```

- a. Execute the pg_demo.sql file using psql to Create the database and add the airplane table
psql -h localhost -p 5432 -U postgres -f pg_demo.sql
 - b. Postgres SQL quick reference - <https://www.postgresqltutorial.com/postgresql-cheat-sheet/>
- 8) Write a Java application to read (select) and write (insert) from Postgres
- a. Install IntelliJ IDE on the Ubuntu Linux VM
 - b. Install openjdk 11 on Ubuntu VM:
<https://www.digitalocean.com/community/tutorials/how-to-install-java-with-apt-on-ubuntu-20-04>
 - c. Install Build tool, Gradle, on Ubuntu VM using snap:
<https://snapcraft.io/install/gradle/ubuntu>
 - d. Gradle Configuration - build.gradle
 - a. Add the following to the top of the build.gradle file to configure plugins, and add a dependency to have Gradle install the Postgres Java client library

```
plugins{
    //Apply the java plugin to add support for Java
    id 'java'

    //Apply the application plugin to add support for building a CLI application.
    id 'application'

    //Apply the IntelliJ IDEA gradle plugin
    id 'idea'
}

// In this section you declare where to find the dependencies of your project
repositories {
    // Use jcenter for resolving your dependencies.
    // You can declare any Maven/Ivy/file repository here.
    jcenter()
}

dependencies {
    // https://mvnrepository.com/artifact/org.postgresql/postgresql
    // Installs the Java JDBC client for Postgres
    // Java Database Connectivity (JDBC) is an application programming interface (API) for the programming language Java
    // implementation group: 'org.postgresql', name: 'postgresql', version: '42.2.23'
}
```

- b. Add the following to the bottom of the build.gradle file to be able to run the Java application

```
application {
    //Define the main class for the application.
    mainClassName = 'com.suncountry.sqldemo.Main'
}
```

- e. Write Java Class (Pg.java) to connect to the Postgres Database, insert rows into the airplane table, and then perform a select to retrieve rows from the airplane table.
- a. Postgres Java Client Library documentation:
https://www.tutorialspoint.com/postgresql/postgresql_java.htm
 - b. Setup the driver for Postgres and connect to the Postgres database (pg_testdb)

```
Class.forName("org.postgresql.Driver");
```

```

        connection = DriverManager.getConnection("jdbc:postgresql://localhost:5432/pg_testdb",
"postgres", "pw");
        connection.setAutoCommit(false);
        System.out.println("Opened postgres database successfully");

```

- c. Select rows from Postgres - Code snippet for selecting rows into a result set "rs" from the airplane table

```

ResultSets =statement.executeQuery( "SELECT * FROM airplane;");
while( rs.next() ) {
    int airplaneid =rs.getInt("airplaneid");
    String callsign =rs.getString("callsign");
    String aircraft_type =rs.getString("aircraft_type");
    int passenger_capacity =rs.getInt("passenger_capacity");
    int max_altitude =rs.getInt("max_altitude");
    int max_speed =rs.getInt("max_speed");
    System.out.println( "airplaneid = "+airplaneid );
    System.out.println( "callsign = "+callsign );
    System.out.println( "aircraft_type = "+aircraft_type );
    System.out.println( "passenger_capacity = "+passenger_capacity );
    System.out.println( "max_altitude = "+max_altitude );
    System.out.println( "max_speed = "+max_speed );
    System.out.println();
}
rs.close();

```

- d. Main to invoke the insert and select methods

9) Install a Microsoft SQL Server Docker Container, run it, and connect to it with sqlcmd CLI

- a. Get the MS SQL Server image from Docker Hub
https://hub.docker.com/_/microsoft-mssql-server
`sudo docker pull mcr.microsoft.com/mssql/server:2019-CU12-ubuntu-20.04`
- b. Start the MS SQL Server Container
`sudo docker run -e "ACCEPT_EULA=Y" -e "SA_PASSWORD=Pw123456*" --name mssql -p 1433:1433 -d mcr.microsoft.com/mssql/server:2019-CU12-ubuntu-20.04`
- c. Connect to the MS SQL Server running in the docker container using CLI tools
 - a. Install sqlcmd - Command-line query tool for MS SQL Server
[Install SQL Server command-line tools on Linux - SQL Server | Microsoft Docs](#)
 - 1) `sudo apt install msodbcsql1`
 - 2) `sudo apt install libodbc1:amd64`
 - b. `sqlcmd help`
[Working with the SQL Server command line \(sqlcmd\) \(sqlshack.com\)](#)
 - c. Connect to a MS SQL Server interactive session
`sqlcmd -S localhost -U SA -P Pw123456*`
 - d. Use sqlcmd to run a script in the mssql_demo.sql file to create the sql_demo database and airplane table
`sqlcmd -S localhost -U SA -P Pw123456* -i mssql_demo.sql`
- d. Create the database using T-SQL (Transact SQL) --- This is what Microsoft calls SQL
 - a. Create a database

```

USE master;
GO
IF DB_ID ( N'Music' ) IS NOT NULL
DROP DATABASE sql_demo;
GO
CREATE DATABASE sql_demo;
GO

```

- 1) List databases:
`sp_databases`
`go`

- b. Create the airplane table in MS SQL Server - Edit a file named `mssql_demo.sql`. This will create a table called `airplane` and insert 2 rows

```
USE sql_demo
IF OBJECT_ID ('dbo.airplane', 'U') IS NOT NULL
    DROP TABLE airplane;
GO
CREATE TABLE airplane
(
    airplaneid int,
    callsign varchar (50),
    aircraft_type varchar(50),
    passenger_capacity int,
    max_altitude int,
    max_speed int
);
GO

INSERT airplane
    (airplaneid, callsign, aircraft_type, passenger_capacity, max_altitude, max_speed)
VALUES
    (12345678, 'SC737', '737', 354, 40000, 540);

INSERT airplane
    (airplaneid, callsign, aircraft_type, passenger_capacity, max_altitude, max_speed)
VALUES
    (12345679, 'SC747', '747', 600, 47000, 640);
GO
```

- 1) Run the SQL file
`sqlcmd -S localhost -U SA -P Pw123456* -i mssql_demo.sql`
- 2) View the table
`use sql_demo`
`select * from airplane`
`go`

- 10) Write a Java Class (`MsSQL.java`) that contains methods to read (select) and write (insert) from MS SQL Server

- a. [Java Ubuntu \(microsoft.com\)](https://docs.microsoft.com/en-us/java/javase/8/platform/connect/quickstart)
- b. Configure the Java client for MS SQL by setting the `Class.forName` and then set the connection to the database (`sql_demo`).

```
Class.forName("com.microsoft.sqlserver.jdbc.SQLServerDriver");
connection = DriverManager
    .getConnection("jdbc:sqlserver://localhost;database=sql_demo;user=sa;password=Pw123456*");
connection.setAutoCommit(false);
```

- c. Add dependency for the Microsoft JDBC driver to Gradle to fix the error:
`java.lang.ClassNotFoundException: com.microsoft.sqlserver.jdbc.SQLServerDriver`. Add to the dependencies section in `build.gradle`

```
// https://mvnrepository.com/artifact/com.microsoft.sqlserver/mssql-jdbc
implementation group: 'com.microsoft.sqlserver', name: 'mssql-jdbc', version: '9.4.0.jre11'
```

- 11) Run the Java Application that reads rows from Microsoft SQL Server, transforms the data, and then inserts the rows into the Postgres database.
- a. From a terminal in the project directory
 - a. `./gradlew run`
 - b. Or run it from within IntelliJ IDEA